

DeepHack.CISS



Convolutional Neural Networks (for NLP)

Valentin Malykh

Overview of today


- From RNNs to CNNs
- CNN Variant1: Simple single layer
- Application: Sentence classification
- More details and tricks
- Evaluation
- Comparison between sentence models
- CNN Variant2: Complex multi layer

From RNNs to CNNs




From RNNs to CNNs

- Recursive neural nets require a parser to get tree structure



- Recurrent neural nets cannot capture phrases without prefix context
And often capture too much of last words in final vector



From RNNs to CNNs

- RNN: Get compositional vectors for grammatical phrases only
- CNN: What if we compute vectors for every possible phrase?
- Example: “the country of my birth” computes vectors for:
 - the country, country of, of my, my birth, the country of, Country of my, of my birth, the country of my, country of my birth
- Regardless of whether it is grammatical
- Wouldn't need parser
- Not very linguistically or cognitively plausible

What is convolution anyway?

- 1d discrete convolution generally:
$$(f * g)[n] = \sum_{m=-M}^M f[n - m]g[m].$$
- Convolution is great to extract features from images

- 2d example
- Yellow shows filter weights
- Green shows input

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 _{x1} | 1 _{x0} | 1 _{x1} | 0 | 0 |
| 0 _{x0} | 1 _{x1} | 1 _{x0} | 1 | 0 |
| 0 _{x1} | 0 _{x0} | 1 _{x1} | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |


Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

From RNNs to CNNs

- First layer: compute all bigram vectors



- Same computation as in RNN but for every pair


$$p = \tanh \left(W \begin{bmatrix} c_1 \\ c_2 \end{bmatrix} + b \right)$$

- This can be interpreted as a convolution over the word vectors

From RNNs to CNNs


- Now multiple options to compute higher layers.
- First option (simple to understand but not necessarily best)
- Just repeat with different weights:

$$p = \tanh \left(W^{(2)} \begin{bmatrix} c_1 \\ c_2 \end{bmatrix} + b \right)$$




From RNNs to CNNs

- First option (simple to understand but not necessarily best)




From RNNs to CNNs

- First option (simple to understand but not necessarily best)




Single Layer CNN

- A simple variant using one convolutional layer and **pooling**
- Based on Collobert and Weston (2011) and Kim (2014)
“Convolutional Neural Networks for Sentence Classification”
- Word vectors: $\mathbf{x}_i \in \mathbb{R}^k$
- Sentence: $\mathbf{x}_{1:n} = \mathbf{x}_1 \oplus \mathbf{x}_2 \oplus \dots \oplus \mathbf{x}_n$ (vectors concatenated)
- Concatenation of words in range: $\mathbf{x}_{i:i+j}$
- Convolutional filter: $\mathbf{w} \in \mathbb{R}^{hk}$ (goes over window of h words)
- Could be 2 (as before) higher, e.g. 3:




Single layer CNN

- Filter w is applied to all possible windows (concatenated vectors)
- Sentence: $\mathbf{x}_{1:n} = \mathbf{x}_1 \oplus \mathbf{x}_2 \oplus \dots \oplus \mathbf{x}_n$
- All possible windows of length h : $\{\mathbf{x}_{1:h}, \mathbf{x}_{2:h+1}, \dots, \mathbf{x}_{n-h+1:n}\}$
- Result is a feature map: $\mathbf{c} = [c_1, c_2, \dots, c_{n-h+1}] \in \mathbb{R}^{n-h+1}$



Single layer CNN

- Filter w is applied to all possible windows (concatenated vectors)
- Sentence: $\mathbf{x}_{1:n} = \mathbf{x}_1 \oplus \mathbf{x}_2 \oplus \dots \oplus \mathbf{x}_n$
- All possible windows of length h : $\{\mathbf{x}_{1:h}, \mathbf{x}_{2:h+1}, \dots, \mathbf{x}_{n-h+1:n}\}$
- Result is a feature map: $\mathbf{c} = [c_1, c_2, \dots, c_{n-h+1}] \in \mathbb{R}^{n-h+1}$



Single layer CNN: Pooling layer

- New building block: Pooling
- In particular: max-over-time pooling layer
- Idea: capture most important activation (maximum over time)
- From feature map $\mathbf{c} = [c_1, c_2, \dots, c_{n-h+1}] \in \mathbb{R}^{n-h+1}$
- Pooled single number: $\hat{c} = \max\{\mathbf{c}\}$
- But we want more features!

Solution: Multiple filters

- Use multiple filter weights w
- Useful to have different window sizes h
- Because of max pooling $\hat{c} = \max\{\mathbf{c}\}$, length of \mathbf{c} irrelevant
$$\mathbf{c} = [c_1, c_2, \dots, c_{n-h+1}] \in \mathbb{R}^{n-h+1}$$
- So we can have some filters that look at unigrams, bigrams, trigrams, 4-grams, etc.

Multi-channel idea

- Initialize with pre-trained word vectors (e.g. word2vec)
- Start with two copies
- Backprop into only one set, keep other “static”
- Both channels are added to c before max-pooling

Classification after one CNN layer

- First one convolution, followed by one max-pooling
- To obtain final feature vector: $\mathbf{z} = [\hat{c}_1, \dots, \hat{c}_m]$
(assuming m filters w)
- Simple final softmax layer $y = \text{softmax} \left(W^{(S)} z + b \right)$

Figure from Kim (2014)




Figure 1: Model architecture with two channels for an example sentence.

n words (possibly zero padded) and each word vector has k dimensions

Tricks to make it work better: Dropout

- Idea: randomly mask/dropout/set to 0 some of the feature weights z
- Create masking vector r of Bernoulli random variables with probability p (a hyperparameter) of being 1

- Delete features during training:

$$y = \textit{softmax} \left(W^{(S)}(r \circ z) + b \right)$$

- Reasoning: Prevents co-adaptation (overfitting to seeing specific feature constellations)

Tricks to make it work better: Dropout

- Idea: randomly mask/dropout/set to 0 some of the feature weights z
- Create masking vector r of Bernoulli random variables with probability p (a hyperparameter) of being 1

- Delete features during training:

$$y = \textit{softmax} \left(W^{(S)}(r \circ z) + b \right)$$

- Reasoning: Prevents co-adaptation (overfitting to seeing specific feature constellations)

Tricks to make it work better: Dropout

$$y = \text{softmax} \left(W^{(S)}(r \circ z) + b \right)$$

- At training time, gradients are backpropagated only through those elements of z vector for which r is 1
- At test time, there is no dropout, so feature vectors z are
- Hence, we scale final vector by Bernoulli probability p

$$\hat{W}^{(S)} = pW^{(S)}$$

- Kim (2014) reports **2-4% improved accuracy** and ability to use very large networks without overfitting

Another regularization trick

- Somewhat less common
- Constrain ℓ_2 norms of weight vectors of each class (row in softmax weight $W_{(s)}$) to fixed number s (also a hyperparameter)
- If $\|W_{c \cdot}^{(S)}\| > s$, then rescale it so $\|W_{c \cdot}^{(S)}\| = s$

All hyperparameters in Kim (2014)

- Find hyperparameters based on dev set
- Nonlinearity: reLu
- Window filter sizes $h = 3, 4, 5$
- Each filter size has 100 featuremaps
- Dropout $p = 0.5$
- L2 constraint s for rows of softmax $s = 3$
- Mini batch size for SGD training: 50
- Word vectors: pre-trained with word2vec, $k = 300$
- During training, keep checking performance on dev set and pick highest accuracy weights for final evaluation

Experiments

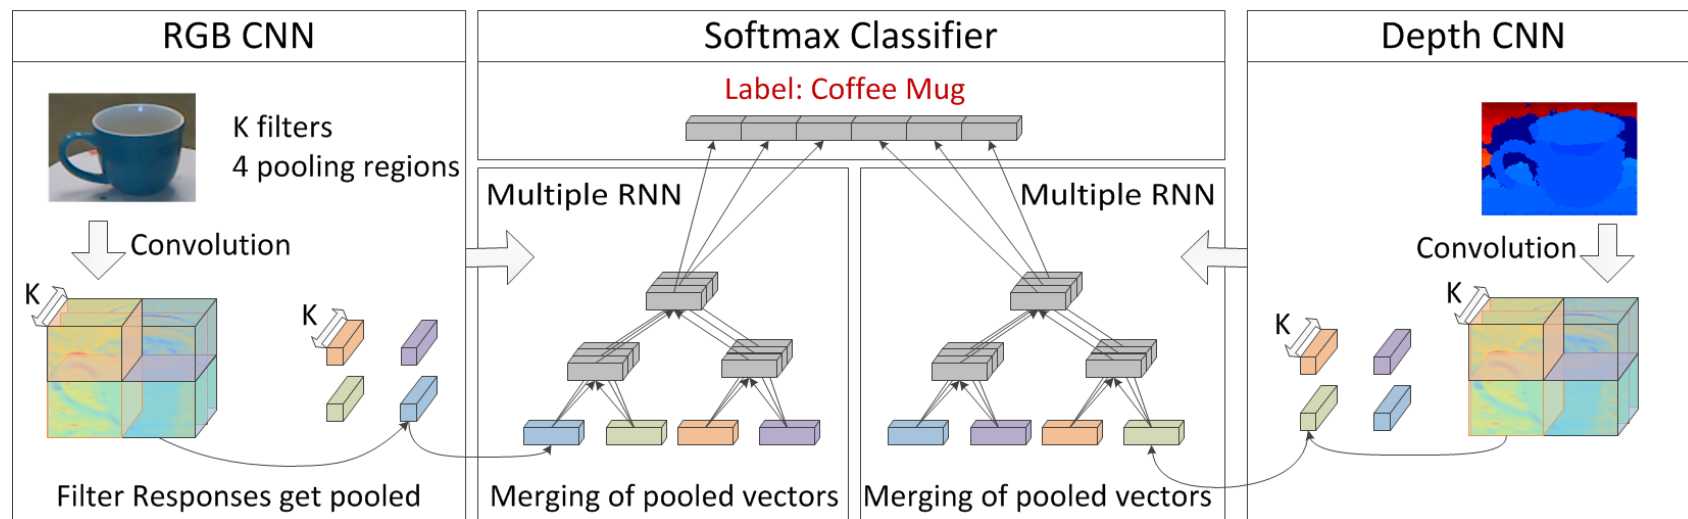
| Model | MR | SST-1 | SST-2 | Subj | TREC | CR | MPQA |
|--------------------------------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| CNN-rand | 76.1 | 45.0 | 82.7 | 89.6 | 91.2 | 79.8 | 83.4 |
| CNN-static | 81.0 | 45.5 | 86.8 | 93.0 | 92.8 | 84.7 | 89.6 |
| CNN-non-static | 81.5 | 48.0 | 87.2 | 93.4 | 93.6 | 84.3 | 89.5 |
| CNN-multichannel | 81.1 | 47.4 | 88.1 | 93.2 | 92.2 | 85.0 | 89.4 |
| RAE (Socher et al., 2011) | 77.7 | 43.2 | 82.4 | | 93.0 | | 86.4 |
| MV-RNN (Socher et al., 2012) | 79.0 | 44.4 | 82.9 | | | | |
| RNTN (Socher et al., 2013) | | 45.7 | 85.4 | | | | |
| DCNN (Kalchbrenner et al., 2014) | | 48.5 | 86.8 | | | | |
| Paragraph-Vec (Le and Mikolov, 2014) | | 48.7 | 87.8 | | | | |
| CCAE (Hermann and Blunsom, 2013) | 77.8 | | | | | | 87.2 |
| Sent-Parser (Dong et al., 2014) | 79.5 | | | | | | 86.3 |
| NBSVM (Wang and Manning, 2012) | 79.4 | | | 93.2 | | 81.8 | 86.3 |
| MNB (Wang and Manning, 2012) | 79.0 | | | 93.6 | | 80.0 | 86.3 |
| G-Dropout (Wang and Manning, 2013) | 79.0 | | | 93.4 | | 82.1 | 86.1 |
| F-Dropout (Wang and Manning, 2013) | 79.1 | | | 93.6 | | 81.9 | 86.3 |
| Tree-CRF (Nakagawa et al., 2010) | 77.3 | | | | | 81.4 | 86.1 |
| CRF-PR (Yang and Cardie, 2014) | | | | | | 82.7 | |
| SVMs (Silva et al., 2011) | | | | | 95.0 | | |

Table 2: Results of our CNN models against other methods. **RAE**: Recursive Autoencoders with pre-trained word vectors from

Problem with comparison?

- Dropout gives 2 – 4 % accuracy improvement
- Several baselines didn't use dropout
- Still remarkable results and simple architecture!
- Difference to window and RNN architectures we described in previous lectures: pooling, many filters and dropout
- Ideas can be used in RNNs too
- Tree-LSTMs obtain better performance on sentence datasets


- Fixed tree RNNs explored in computer vision:
Socher et al (2012): “Convolutional-Recursive Deep Learning for 3D Object Classification”




Relationship between RNNs and CNNs

-

CNN




RNN




Relationship between RNNs and CNNs



CNN




RNN




Relationship between RNNs and CNNs

-

CNN




RNN




- **Stride size** flexible in CNNs, RNNs “weighted average pool”
- Tying (sharing) weights of filters inside vs across different layers
- CNN: multiple filters, additional layer type: max-pooling
- Balanced input independent structure vs input specific tree

CNN alternatives

- Narrow vs wide convolution




- Complex pooling schemes (over sequences) and deeper convolutional layers
- Kalchbrenner et al. (2014)



CNN application: Translation

- One of the first successful neural machine translation efforts
- Uses CNN for encoding and RNN for decoding
- Kalchbrenner and Blunsom (2013)
“Recurrent Continuous Translation Models”




Model comparison

- **Bag of Vectors:** Surprisingly good baseline for simple classification problems. Especially if followed by a few layers!
- **Window Model:** Good for single word classification for problems that do not need wide context
- **CNNs:** good for classification, unclear how to incorporate phrase level annotation (can only take a single label), need zero padding for shorter phrases, hard to interpret, easy to parallelize on GPUs

Model comparison

- **Recursive Neural Networks:** most linguistically plausible, interpretable, provide most important phrases (for visualization), need parse trees
- **Recurrent Neural Networks:** Most cognitively plausible (reading from left to right), not usually the highest classification performance but lots of improvements right now with gates (GRUs, LSTMs, etc).
- Best but also most complex models: Hierarchical recurrent neural networks with attention mechanisms and additional memory

Quasi-Recurrent Neural Network



- Parallelism computation across time:

$$\mathbf{z}_t = \tanh(\mathbf{W}_z^1 \mathbf{x}_{t-1} + \mathbf{W}_z^2 \mathbf{x}_t)$$

$$\mathbf{f}_t = \sigma(\mathbf{W}_f^1 \mathbf{x}_{t-1} + \mathbf{W}_f^2 \mathbf{x}_t)$$

$$\mathbf{o}_t = \sigma(\mathbf{W}_o^1 \mathbf{x}_{t-1} + \mathbf{W}_o^2 \mathbf{x}_t).$$

$$\mathbf{Z} = \tanh(\mathbf{W}_z * \mathbf{X})$$

$$\mathbf{F} = \sigma(\mathbf{W}_f * \mathbf{X})$$

$$\mathbf{O} = \sigma(\mathbf{W}_o * \mathbf{X}),$$

- Element-wise gated recurrence for parallelism across channels:


$$\mathbf{h}_t = \mathbf{f}_t \odot \mathbf{h}_{t-1} + (1 - \mathbf{f}_t) \odot \mathbf{z}_t,$$

Q-RNNs for Language Modeling

- Better

| Model | Parameters | Validation | Test |
|---|------------|------------|------|
| LSTM (medium) (Zaremba et al., 2014) | 20M | 86.2 | 82.7 |
| Variational LSTM (medium) (Gal & Ghahramani, 2016) | 20M | 81.9 | 79.7 |
| LSTM with CharCNN embeddings (Kim et al., 2016) | 19M | — | 78.9 |
| Zoneout + Variational LSTM (medium) (Merity et al., 2016) | 20M | 84.4 | 80.6 |
| <i>Our models</i> | | | |
| LSTM (medium) | 20M | 85.7 | 82.0 |
| QRNN (medium) | 18M | 82.9 | 79.9 |
| QRNN + zoneout ($p = 0.1$) (medium) | 18M | 82.1 | 78.3 |

- Faster



| | | Sequence length | | | | |
|------------|-----|-----------------|------|-------|-------|-------|
| | | 32 | 64 | 128 | 256 | 512 |
| Batch size | 8 | 5.5x | 8.8x | 11.0x | 12.4x | 16.9x |
| | 16 | 5.5x | 6.7x | 7.8x | 8.3x | 10.8x |
| | 32 | 4.2x | 4.5x | 4.9x | 4.9x | 6.4x |
| | 64 | 3.0x | 3.0x | 3.0x | 3.0x | 3.7x |
| | 128 | 2.1x | 1.9x | 2.0x | 2.0x | 2.4x |
| | 256 | 1.4x | 1.4x | 1.3x | 1.3x | 1.3x |

Q-RNNs for Sentiment Analysis

- Often better and faster than LSTMs

| Model | Time / Epoch (s) | Test Acc (%) |
|--|------------------|--------------|
| BSVM-bi (Wang & Manning, 2012) | — | 91.2 |
| 2 layer sequential BoW CNN (Johnson & Zhang, 2014) | — | 92.3 |
| Ensemble of RNNs and NB-SVM (Mesnil et al., 2014) | — | 92.6 |
| 2-layer LSTM (Longpre et al., 2016) | — | 87.6 |
| Residual 2-layer bi-LSTM (Longpre et al., 2016) | — | 90.1 |
| <i>Our models</i> | | |
| Deeply connected 4-layer LSTM (cuDNN optimized) | 480 | 90.9 |
| Deeply connected 4-layer QRNN | 150 | 91.4 |
| D.C. 4-layer QRNN with $k = 4$ | 160 | 91.1 |

- More interpretable

- Example:

- Initial positive review

- *Review starts out positive*

At 117: “not exactly a bad story”

At 158: “I recommend this movie to everyone, even if you’ve never played the game”

